Inference from a CT volume

64 distinct patients · A100 80 GB · 12 vCPU · batch one

Bar = mean · dot = p95 · annotation = mean / p95 seconds

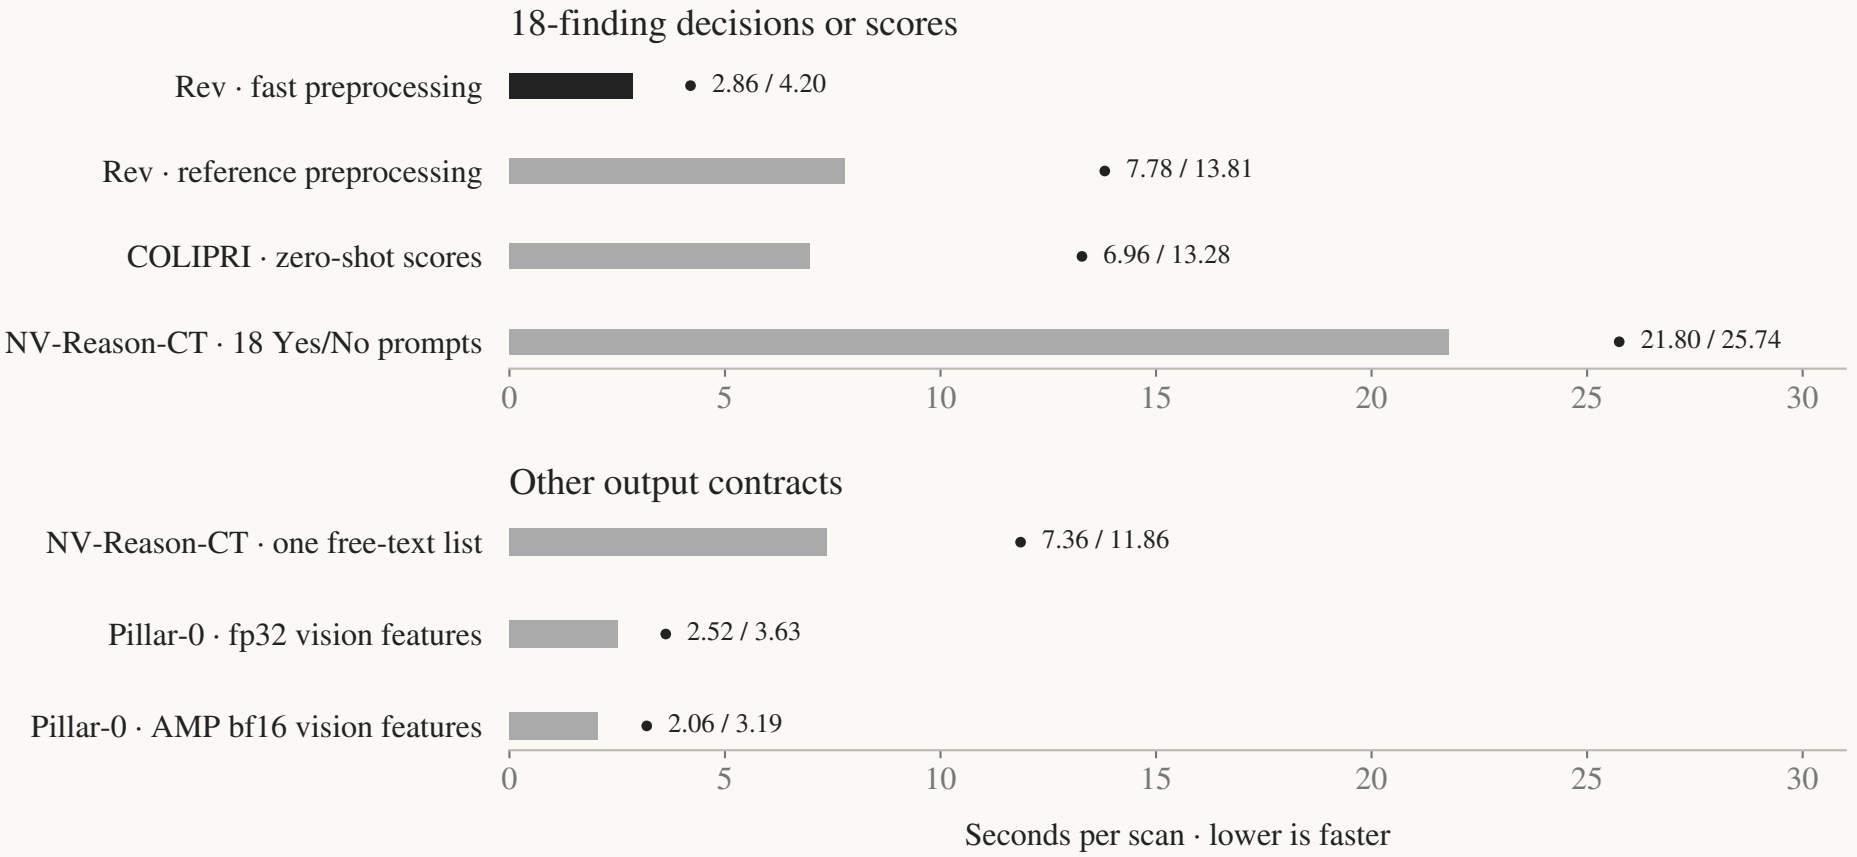

Warm local NIfTI → CPU-visible output; downloads and model loading excluded. Same scans in all rows.
Rev uses fp32; NV and COLIPRI use bf16. Pillar rows produce features, with no trained finding head.
The NV list protocol produced 0/64 requested 18-key JSON mappings. See the benchmark report for protocols.